

reasoning about choreographic programs

luís cruz-filipe

with eva graversen, fabrizio montesi & marco peressotti

department of mathematics and computer science
university of southern denmark

concurrency and logic seminar
june 27th, 2023

choreographic programming, conceptually

what are choreographies?

high-level global specifications of concurrent and distributed systems

a programming paradigm with good properties

implementations for the local endpoints are automatically generated

- guaranteed to be deadlock-free
- guaranteed to satisfy the specification

an example

authentication choreography

```
X = c.credentials --> ip.x;  
  If ip.(check x)  
    Then ip --> c[left]; ip.go --> s.b; s.token --> c.t  
    Else ip --> c[right]; X
```

an example

authentication choreography

```
X = c.credentials --> ip.x;  
  If ip.(check x)  
    Then ip --> c[left]; ip.go --> s.b; s.token --> c.t  
    Else ip --> c[right]; X
```

the problem

how can we reason formally about what this choreography does?

- currently: *ad-hoc* properties can be proved by simulating execution
- better: have a more general framework for reasoning about choreographies

our contribution

we propose a sound and complete hoare calculus for a simple choreography language

our contribution

we propose a sound and complete hoare calculus for a simple choreography language

design choices

avoid *ad-hoc* solutions

- use a pre-existing choreography language
 ↪ for the good and for the bad
- design a “standard” hoare calculus
- separation of concerns
- emphasis on parameterisation

a simple choreography language

syntax

choreography bodies are defined by the following grammar

$$C ::= I; C \mid \text{if } p.b \text{ then } C_1 \text{ else } C_2 \mid X \mid [\vec{q}, X] C \mid \mathbf{0}$$

$$I ::= p.x := e \mid p.e \rightarrow q.x \mid p \rightarrow q[L]$$

- $p.x := e$: local computation
- $p.e \rightarrow q.x$: value communication
- $p \rightarrow q[L]$: label selection
- X : procedure call
- $[\vec{q}, X] C$: runtime term for partially entered procedures

a simple choreography language

syntax

choreography bodies are defined by the following grammar

$$\begin{aligned}
 C &::= I; C \mid \text{if } p.b \text{ then } C_1 \text{ else } C_2 \mid X \mid \lceil \vec{q}, X \rceil C \mid \mathbf{0} \\
 I &::= p.x := e \mid p.e \rightarrow q.x \mid p \rightarrow q[L]
 \end{aligned}$$

choreographic programs

$\langle \mathcal{C}, C \rangle$ where \mathcal{C} maps procedure names to their definitions

semantics

ltl semantics

configurations are pairs $\langle C, \Sigma \rangle$ where Σ is a memory state

- $\Sigma(p)(x)$ returns the value stored at p 's variable x
- $e \downarrow_{\Sigma(p)} v$ returns the result of locally evaluating e at p using Σ

semantics

ltl semantics

configurations are pairs $\langle C, \Sigma \rangle$ where Σ is a memory state

the rules

three groups of rules

- formalisation of the intuition behind the constructs, e.g. if $e \downarrow_{\Sigma(p)} v$, then

$$\langle p.x := e; C, \Sigma \rangle \xrightarrow{\tau @ p}_{\mathcal{C}} \langle C, \Sigma[\langle p, x \rangle \mapsto v] \rangle$$

semantics

ltl semantics

configurations are pairs $\langle C, \Sigma \rangle$ where Σ is a memory state

the rules

three groups of rules

- formalisation of the intuition behind the constructs
- formalisation of out-of-order execution, e.g. if $\langle C, \Sigma \rangle \xrightarrow{\mu}_{\mathcal{C}} \langle C', \Sigma' \rangle$ and I does not involve processes appearing in μ , then

$$\langle I; C, \Sigma \rangle \xrightarrow{\mu}_{\mathcal{C}} \langle I; C', \Sigma' \rangle$$

semantics

ltl semantics

configurations are pairs $\langle C, \Sigma \rangle$ where Σ is a memory state

the rules

three groups of rules

- formalisation of the intuition behind the constructs
- formalisation of out-of-order execution
- rules allowing processes to enter procedure calls independently
 \rightsquigarrow use runtime terms

the intuition

general idea

- judgements are triples $\{\varphi\}C\{\psi\}$
if C is executed from a state where φ holds and execution terminates, then ψ holds in the final state
- main inference rules match the rules of semantics

the intuition

general idea

- judgements are triples $\{\varphi\}C\{\psi\}$
if C is executed from a state where φ holds and execution terminates, then ψ holds in the final state
- main inference rules match the rules of semantics

three challenges

- what formulas can we write about states?

the intuition

general idea

- judgements are triples $\{\varphi\}C\{\psi\}$
if C is executed from a state where φ holds and execution terminates, then ψ holds in the final state
- main inference rules match the rules of semantics

three challenges

- what formulas can we write about states?
- the usual rule for assignment

$$\frac{}{\{\varphi'\}p.x := e; \mathbf{0}\{\varphi\}}$$

where φ' is obtained from φ by replacing $p.x$ with e does not work

the intuition

general idea

- judgements are triples $\{\varphi\}C\{\psi\}$
if C is executed from a state where φ holds and execution terminates, then ψ holds in the final state
- main inference rules match the rules of semantics

three challenges

- what formulas can we write about states?
- the usual rule for assignment does not work
- how do we deal with procedure calls?

the state logic

an equational logic

- parameterised on the language of expressions in the choreography language
- variables are localised, e.g. $p.x$
- parameterised on a decidable theory \mathfrak{D} whose terms include logical variables \mathcal{X}

the state logic

an equational logic

- parameterised on the language of expressions in the choreography language
- variables are localised, e.g. $p.x$
- parameterised on a decidable theory \mathfrak{D} whose terms include logical variables \mathcal{X}

syntax

$$\varphi ::= (\mathcal{E} = \mathcal{X}) \mid \delta \mid \varphi \wedge \varphi \mid \neg \varphi$$

the state logic

syntax

$$\varphi ::= (\mathcal{E} = \mathcal{X}) \mid \delta \mid \varphi \wedge \varphi \mid \neg \varphi$$

semantics

given an assignment ρ assigning values to logical variables:

$$\frac{\mathcal{E} \downarrow_{\Sigma} \rho(\mathcal{X})}{\Sigma \Vdash_{\rho} \mathcal{E} = \mathcal{X}} \quad \frac{\delta \in \mathfrak{D} \quad \delta \rho \text{ is true}}{\Sigma \Vdash_{\rho} \delta}$$

localisation

definition

- $L(p, e)$ is the logical expression obtained from e by replacing every choreography variable x with $p.x$
- $\mathcal{E}[q.x := p.e]$ is the expression obtained from \mathcal{E} by replacing every occurrence of $q.x$ with $L(p, e)$
- localised substitution extends to formulas in the natural way

localisation

definition

- $L(p, e)$ is the logical expression obtained from e by replacing every choreography variable x with $p.x$
- $\mathcal{E}[q.x := p.e]$ is the expression obtained from \mathcal{E} by replacing every occurrence of $q.x$ with $L(p, e)$
- localised substitution extends to formulas in the natural way

an example

let φ be $p.x = \mathcal{X}$ and e be $y - z$, then:

- $L(p, y - z) = p.y - p.z$
- $\varphi[p.x := p.(y - z)]$ is $p.y - p.z = \mathcal{X}$

localisation

definition

- $L(p, e)$ is the logical expression obtained from e by replacing every choreography variable x with $p.x$
- $\mathcal{E}[q.x := p.e]$ is the expression obtained from \mathcal{E} by replacing every occurrence of $q.x$ with $L(p, e)$
- localised substitution extends to formulas in the natural way

properties

- if $\rho(\mathcal{X}) = v$: then $e \downarrow_{\Sigma(p)} v$ iff $\Sigma \Vdash_{\rho} L(p, e) = \mathcal{X}$
- if $e \downarrow_{\Sigma(p)} v$: then $\Sigma[\langle p, x \rangle \mapsto v] \Vdash_{\rho} \varphi$ iff $\Sigma \Vdash_{\rho} \varphi[q.x := p.e]$ for all ρ

dealing with procedure calls

procedure specification maps

to reason about a program we need a description of the behaviour of the procedures

$$\mathfrak{C}(X) = \langle \varphi_X, \psi_X \rangle$$

- intended meaning: if φ_X holds when X is called and execution terminates, then ψ_X holds in the final state

the rules, part i

about instructions...

$$\frac{\vdash_{\mathcal{C}} \{\varphi\} C\{\psi\}}{\vdash_{\mathcal{C}} \{\varphi[p.x := p.e]\} p.x := e; C\{\psi\}} \text{H|ASSIGN}$$

$$\frac{\vdash_{\mathcal{C}} \{\varphi\} C\{\psi\}}{\vdash_{\mathcal{C}} \{\varphi[q.x := p.e]\} p.e \rightarrow q.x; C\{\psi\}} \text{H|COM}$$

$$\frac{\vdash_{\mathcal{C}} \{\varphi\} C\{\psi\}}{\vdash_{\mathcal{C}} \{\varphi\} p \rightarrow q[L]; C\{\psi\}} \text{H|SEL}$$

*the rules, part ii**... about compound choreographies...*

$$\frac{}{\vdash_{\mathfrak{C}} \{\varphi\} \mathbf{0} \{\varphi\}} \text{H|NIL}$$

$$\frac{\begin{array}{l} \vdash_{\mathfrak{C}} \{\varphi \wedge L(p, b) \stackrel{\mathcal{X}}{=} \text{true}\} C_1 \{\psi\} \\ \vdash_{\mathfrak{C}} \{\varphi \wedge L(p, b) \stackrel{\mathcal{X}}{=} \text{false}\} C_2 \{\psi\} \end{array} \quad \mathcal{X} \text{ fresh}}{\vdash_{\mathfrak{C}} \{\varphi\} \text{if } p.b \text{ then } C_1 \text{ else } C_2 \{\psi\}} \text{H|COND}$$

$$\frac{\mathfrak{C}(X) = \langle \varphi, \psi \rangle}{\vdash_{\mathfrak{C}} \{\varphi\} X \{\psi\}} \text{H|CALL}$$

$$\frac{\vdash_{\mathfrak{C}} \{\varphi\} C \{\psi\}}{\vdash_{\mathfrak{C}} \{\varphi\} [\vec{a}, X] C \{\psi\}} \text{H|CALL'}$$

the rules, part iii

...and structural rules

$$\frac{\mathcal{D} \models \varphi \rightarrow \varphi' \quad \vdash_{\mathcal{C}} \{\varphi'\} C \{\psi'\} \quad \mathcal{D} \models \psi' \rightarrow \psi}{\vdash_{\mathcal{C}} \{\varphi\} C \{\psi\}} \text{H|W}_{\text{EAK}}$$

auxiliary results

head reductions

$\langle C, \Sigma \rangle \xRightarrow[\mathcal{E}]{\mu} \langle C', \Sigma' \rangle$ denotes that C reduces to C' without using out-of-order execution

auxiliary results

head reductions

$\langle C, \Sigma \rangle \xRightarrow[\mathcal{C}]{\mu} \langle C', \Sigma' \rangle$ denotes that C reduces to C' without using out-of-order execution

confluence

choreography execution is confluent

- in particular, if $\langle C, \Sigma \rangle \rightarrow_{\mathcal{C}}^* \langle \mathbf{0}, \Sigma' \rangle$ then also $\langle C, \Sigma \rangle \Rightarrow_{\mathcal{C}}^* \langle \mathbf{0}, \Sigma' \rangle$
- together with determinism, this allows us to focus only on head reductions

soundness

consistency

\mathcal{C} is *consistent* with \mathcal{C} if $\vdash_{\mathcal{C}} \{\varphi_X\} \mathcal{C}(X) \{\psi_X\}$ for every X

soundness

consistency

\mathfrak{C} is *consistent* with \mathcal{C} if $\vdash_{\mathfrak{C}} \{\varphi_X\} \mathcal{C}(X) \{\psi_X\}$ for every X

theorem

- \mathfrak{C} is consistent with \mathcal{C}
 - $\vdash_{\mathfrak{C}} \{\varphi\} C \{\psi\}$
 - $\Sigma \Vdash_{\rho} \varphi$
 - $\langle C, \Sigma \rangle \rightarrow_{\mathcal{C}}^* \langle \mathbf{0}, \Sigma' \rangle$
- } implies $\Sigma' \Vdash_{\rho} \psi$

weakest liberal preconditions

definition

$$\text{wlp}_{\mathcal{C}}((p.x := e; C), \psi) = \text{wlp}_{\mathcal{C}}(C, \psi)[p.x := p.e]$$

$$\text{wlp}_{\mathcal{C}}((p.e \rightarrow q.x; C), \psi) = \text{wlp}_{\mathcal{C}}(C, \psi)[q.x := p.e]$$

$$\text{wlp}_{\mathcal{C}}((p \rightarrow q[L]; C), \psi) = \text{wlp}_{\mathcal{C}}(C, \psi)$$

$$\begin{aligned} \text{wlp}_{\mathcal{C}}(\text{if } p.b \text{ then } C_1 \text{ else } C_2, \psi) = & (L(p, b) \stackrel{\mathcal{X}}{=} \text{true} \rightarrow \text{wlp}_{\mathcal{C}}(C_1, \psi)) \\ & \wedge (L(p, b) \stackrel{\mathcal{X}}{=} \text{false} \rightarrow \text{wlp}_{\mathcal{C}}(C_2, \psi)) \end{aligned}$$

$$\text{wlp}_{\mathcal{C}}(X, \psi) = \varphi_X$$

$$\text{wlp}_{\mathcal{C}}(\lceil \vec{q}, X \rceil C, \psi) = \text{wlp}_{\mathcal{C}}(C, \psi)$$

$$\text{wlp}_{\mathcal{C}}(\mathbf{0}, \psi) = \psi$$

\rightsquigarrow essentially read the rules “backwards”

partial completeness

adequacy

\mathfrak{C} is *adequate for ψ* given \mathcal{C} if, for all X :

- φ_X is equivalent to $\text{wlp}_{\mathfrak{C}}(\mathcal{C}(X), \psi)$
- $\psi_X = \psi$

partial completeness

adequacy

\mathfrak{C} is *adequate* for ψ given \mathcal{C} if, for all X :

- φ_X is equivalent to $\text{wlp}_{\mathfrak{C}}(\mathcal{C}(X), \psi)$
- $\psi_X = \psi$

lemma

if \mathfrak{C} is adequate for ψ given \mathcal{C} , then \mathfrak{C} is consistent with \mathcal{C}

↪ can be combined with soundness

partial completeness

adequacy

\mathfrak{C} is *adequate* for ψ given \mathcal{C} if, for all X :

- φ_X is equivalent to $\text{wlp}_{\mathfrak{C}}(\mathcal{C}(X), \psi)$
- $\psi_X = \psi$

theorem

- \mathfrak{C} is adequate for ψ given \mathcal{C}
 - whenever $\Sigma \Vdash_{\rho} \varphi$ and $\langle C, \Sigma \rangle \rightarrow_{\mathcal{C}}^* \langle \mathbf{0}, \Sigma' \rangle$, then $\Sigma' \Vdash_{\rho} \psi$
- } implies $\vdash_{\mathfrak{C}} \{\varphi\} C \{\psi\}$

(un)decidability

the best...

the judgement $\vdash_{\mathcal{C}} \{\varphi\} C \{\psi\}$ is decidable

(un)decidability

the best...

the judgement $\vdash_{\mathfrak{C}} \{\varphi\} C \{\psi\}$ is decidable

... the good...

if the set of procedure names is finite:

- consistency between \mathfrak{C} and \mathcal{C} is decidable
- adequacy of \mathfrak{C} for ψ and \mathcal{C} is decidable

(un)decidability

the best...

the judgement $\vdash_{\mathfrak{C}} \{\varphi\} C \{\psi\}$ is decidable

... the good...

if the set of procedure names is finite:

- consistency between \mathfrak{C} and \mathcal{C} is decidable
- adequacy of \mathfrak{C} for ψ and \mathcal{C} is decidable

... and the not-so-good

there is no algorithm that, given \mathcal{C} and ψ , always returns \mathfrak{C} that is adequate for ψ given \mathcal{C}

↪ proof idea: $\text{wlp}_{\mathfrak{C}}(X, \perp) = \top$ iff execution of X always diverges

a note on divergence

the big minus

- our calculus only proves properties of terminating executions
- since we do not have sequential composition, the target formula in all judgements holds in the (same) final state

a note on divergence

the big minus

- our calculus only proves properties of terminating executions
- since we do not have sequential composition, the target formula in all judgements holds in the (same) final state

↪ but hey, you gotta start somewhere!

a note on divergence

the big minus

- our calculus only proves properties of terminating executions
- since we do not have sequential composition, the target formula in all judgements holds in the (same) final state

a closer look at consistency

- whenever X is called, φ_X must hold
- because of out-of-order execution, there is no guarantee that the choreography ever passes those points...

a note on divergence

the big minus

- our calculus only proves properties of terminating executions
- since we do not have sequential composition, the target formula in all judgements holds in the (same) final state

a closer look at consistency

- whenever X is called, φ_X must hold
- because of out-of-order execution, there is no guarantee that the choreography ever passes those points...

future work – explore this idea

- formally prove the informal statement above
- capitalise on confluence to get stronger results

in a nutshell

- a hoare calculus for a simple choreography language
- agnostic, modular, generalisable
- soundness, partial completeness and (some) decidability
- potentially extendable to reasoning about non-terminating systems: liveness, reactivity, ...

thank you!